

Self-Supervised Transformers as Iterative Solution Improvers for Constraint Satisfaction



UNIVERSITY OF
TORONTO



ICML
International Conference
On Machine Learning

Yudong Will Xu , Wenhao Li, Elias B. Khalil, Scott Sanner

Motivation

Can Transformers solve Constraint Reasoning Problems?

Existing approaches:

- Supervised Learning: labels are often *difficult to obtain* and sometimes *ambiguous*.
- Reinforcement Learning: uses black box reward function which are often complex and *difficult to train*.

Our approach: Self-Supervised Learning.

Background


Constraint Satisfaction Problem

Variables $X = \{X_{1,1}, X_{1,2}, \dots, X_{9,9}\}$

Domains $D = \{D_1, D_2, \dots, D_{81}\}$
 $D_i = \{1, 2, 3, 4, 5, 6, 7, 8, 9\}$

Constraints $AllDifferent(X_{1,1}, X_{1,2}, \dots, X_{1,9})$
 $AllDifferent(X_{1,1}, X_{2,1}, \dots, X_{9,1})$
 $AllDifferent(X_{1,1}, X_{1,2}, \dots, X_{3,3})$
...

Stochastic Local Search



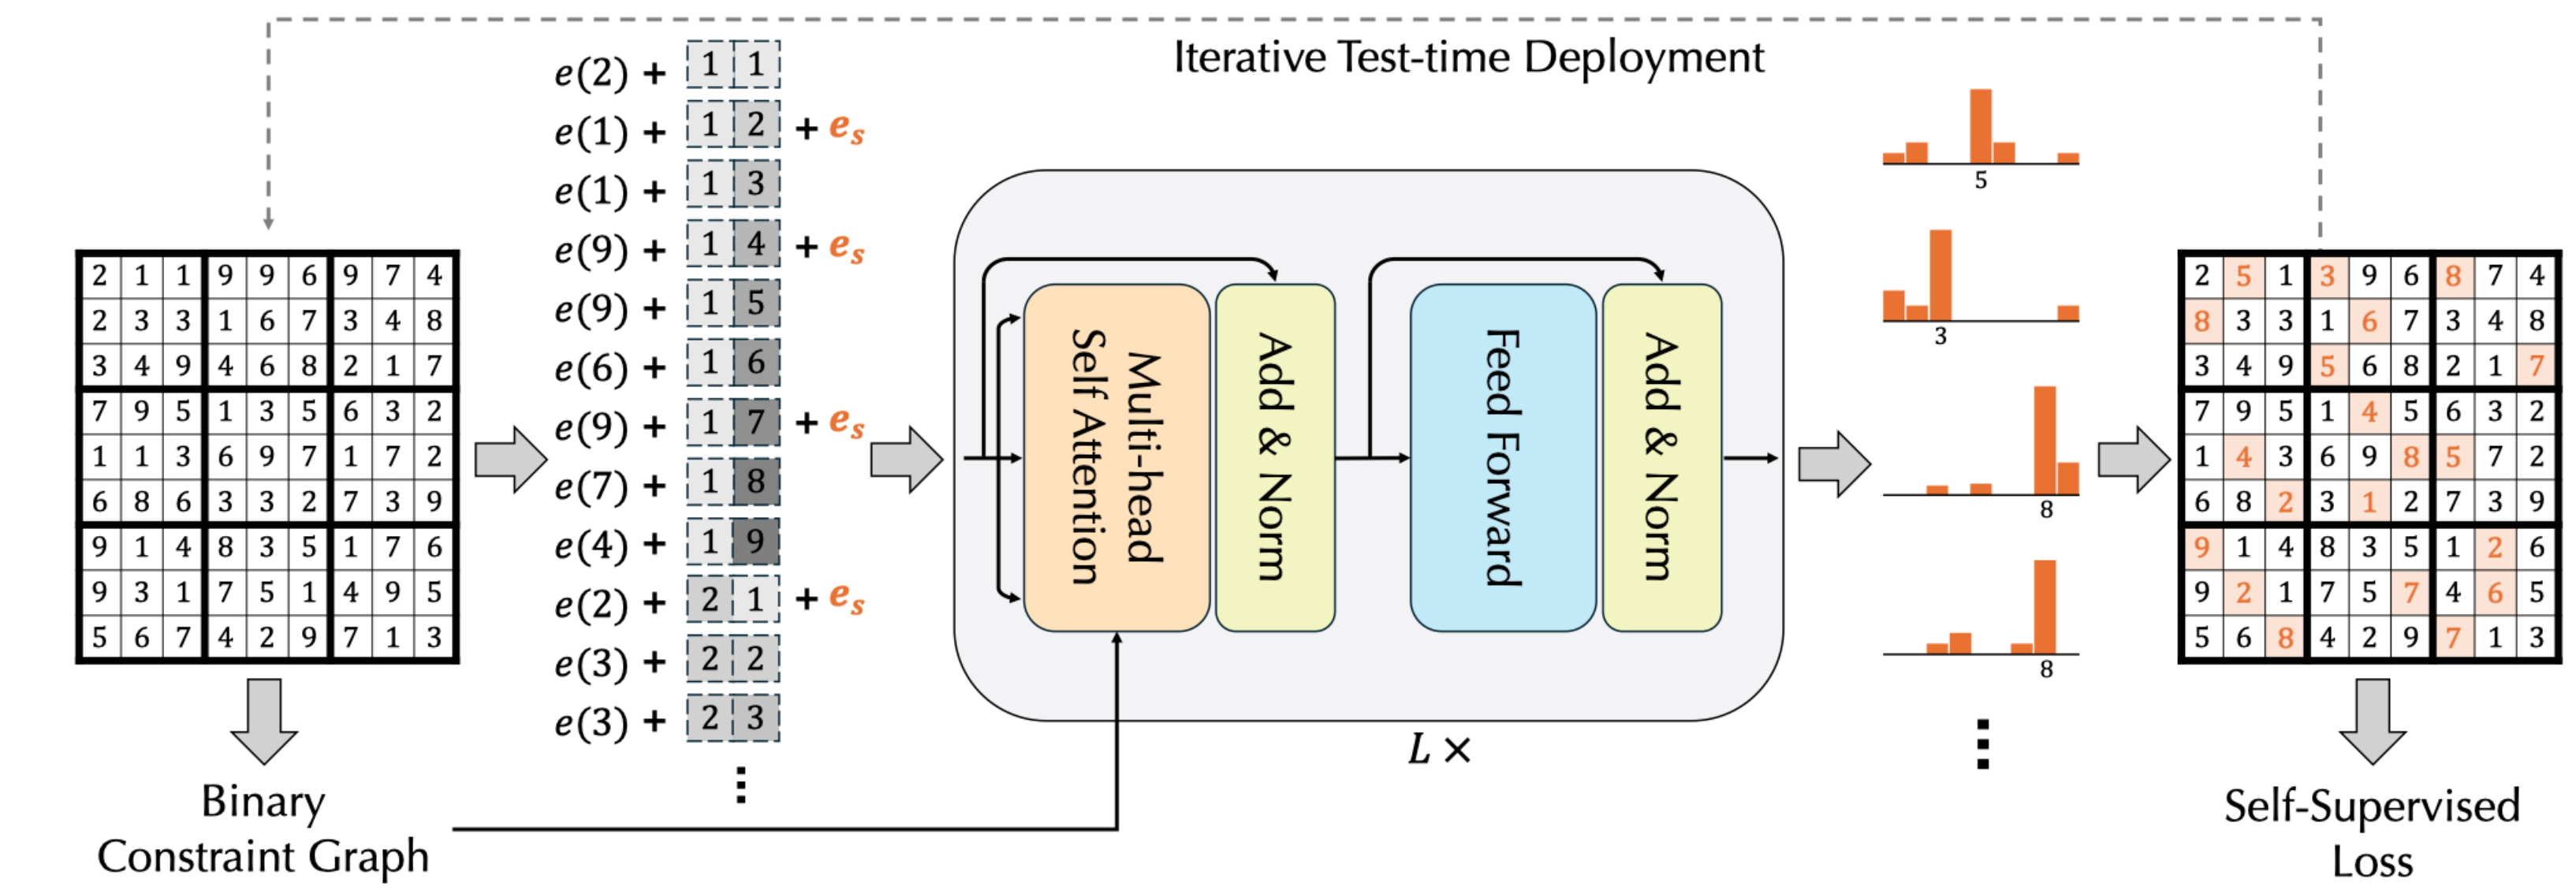
Representation

- Variable indices as absolute positional encodings (APEs)
 $APE(x_{1,2}) = \text{Concat}(\text{PE}(1), \text{PE}(2))$
- Binary constraint graph as relative positional encodings (RPEs)
 $RPE(i, j) = c \cdot \mathbb{I}[(i, j) \notin E]$

Architecture: Single-step Transformer

- Our model mimics local search by randomly selecting a subset of the variables to update, which are then marked by a special subset embedding e_s .
- Variable assignment with Gumbel Softmax.
 $\hat{y}_i = \text{GumbelSoftmax}(\mathbf{W}_{\text{out}} \mathbf{h}_i^{(L)} + \mathbf{b}_{\text{out}})$
 $v'_i = \arg \max \hat{y}_i, \quad \forall i \in S.$

ConsFormer



Training: Self-Supervised Loss

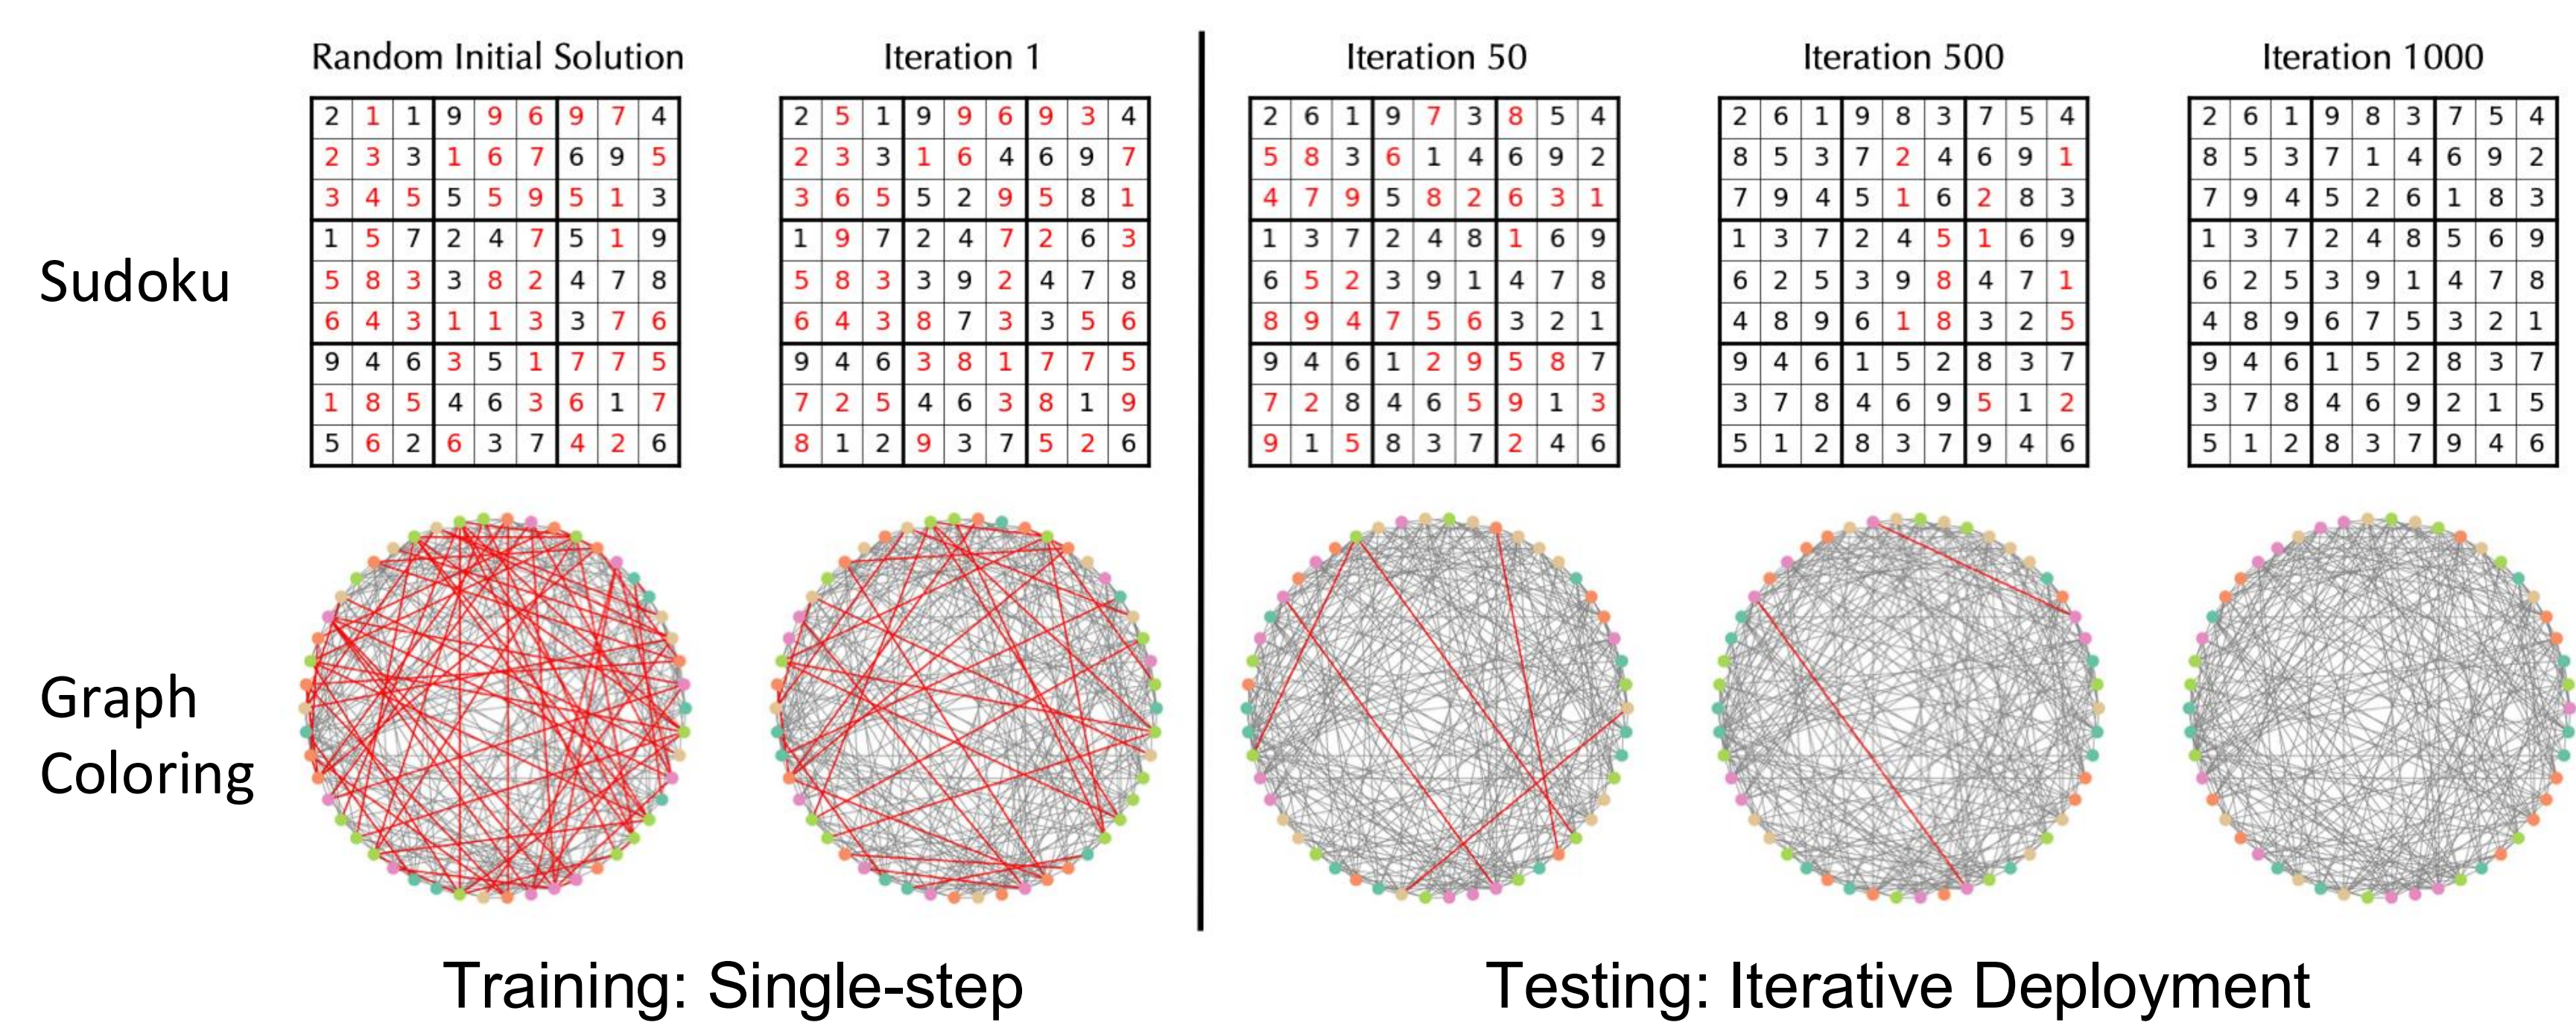
- Define continuous penalty functions p_i for constraint c_i such that
 $p_i(X_i) = 0 \iff c_i(X_i) = \text{True},$
- Loss is then defined by
 $\mathcal{L} = \sum_i \lambda_i f(p_i(X_i)) \quad \forall p_i \in P,$

$ALLDIFFERENT_{m=n}(x_1, \dots, x_n) \implies p = \sum_j |1 - \sum_i x_i^{(j)}|.$

Valid Assignment:
 $x_1 = (1, 0, 0), x_2 = (0, 1, 0), x_3 = (0, 0, 1)$
 $\sum_i x_i^{(1)} = 1, \sum_i x_i^{(2)} = 1, \sum_i x_i^{(3)} = 1$
 $p = |1 - 1| + |1 - 1| + |1 - 1| = 0.$

Invalid Assignment:
 $x_1 = (0.9, 0.1, 0), x_2 = (0.9, 0.1, 0), x_3 = (0, 0, 1)$
 $\sum_i x_i^{(1)} = 1.8, \sum_i x_i^{(2)} = 0.2, \sum_i x_i^{(3)} = 1$
 $p = |1 - 1.8| + |1 - 0.2| + |1 - 1| = 0.8 + 0.8 + 0 = 1.6.$

Experiments



Method	Test Instances	Harder OOD Instances
Wang et al. (2019)	98.3	3.2
Palm et al. (2018)	99.8	28.6
Yang et al. (2023)	100	32.9
Yang et al. (2023) (2k Iters)	97.7	14.0
Du et al. (2024)	99.4	62.1
ConsFormer (2k Iters)	100	65.88
ConsFormer (10k Iters)	100	77.74

Method	Test Instances	Harder OOD Instances
Graph-Coloring-5 ($n = 50 \rightarrow n = 100$)		
OR-Tools (10s)	83.08	57.16
ANYCSP (10s)	79.17	34.83
ConsFormer (10s)	81.00	47.33
Graph-Coloring-10 ($n = 100 \rightarrow n = 200$)		
OR-Tools (10s)	52.41	10.25
ANYCSP (10s)	0.00	0.00
ConsFormer (10s)	52.60	11.92

Key Takeaways

- Iterative single-step reasoning enables out of distribution generalization.
Harder problem \rightarrow Run more iterations.
- Self-supervised Learning enables constraints reasoning *without labeled examples*.

Paper 

Code 

References

Tönshoff, Jan, et al. "One model, any CSP: graph neural networks as fast global search heuristics for constraint satisfaction." *Proceedings of the Thirty-Second International Joint Conference on Artificial Intelligence*, 2023.

Yang, Zhun, Adam Ishay, and Joohyung Lee. "Learning to Solve Constraint Satisfaction Problems with Recurrent Transformer." *The Eleventh International Conference on Learning Representations*, 2023.

Wang, Po-Wei, et al. "Satnet: Bridging deep learning and logical reasoning using a differentiable satisfiability solver." *International Conference on Machine Learning*. PMLR, 2019.

Palm, Rasmus, Ulrich Paquet, and Ole Winther. "Recurrent relational networks." *Advances in neural information processing systems* 31, 2018.

Du, Yilun, Jiayuan Mao, and Joshua B. Tenenbaum. "Learning Iterative Reasoning through Energy Diffusion." *Forty-first International Conference on Machine Learning*, 2024.